



ODDC Critical QA

Quality Assurance Checklist

February 2026 — Confidential

Critical Q&A; ODDC Framework — Tough Questions Answered This document addresses the hardest technical, regulatory, and business objections to runtime enforcement and conformance determination for autonomous systems. These questions are drawn from engagements with regulators, insurers, CTOs, and legal counsel across multiple industries.

Fundamental Objections Q1: Isn't ODDC just adding cost and complexity to systems that are already safe? "Safe" is not a binary condition — it is a claim. Today, that claim is largely unverifiable at runtime. ODDC does not add complexity arbitrarily; it adds accountability where accountability is currently absent. In safety-critical domains, the marginal cost of enforcement is negligible compared to the cost of a failure that cannot be reconstructed, explained, or adjudicated.

We have seen this pattern repeatedly: airbags added cost to vehicles, triple-redundant flight controls added cost to aircraft, cryptographic controls added overhead to financial systems. Each was initially dismissed as excessive.

Each became mandatory once society decided that provable safety mattered more than marginal efficiency.

Q2: Why can't the AI just police itself? Because self-policing systems fail precisely when trust is most needed.

We do not allow pilots to certify their own flights, banks to audit their own books, or software to attest to its own integrity. An AI enforcing its own limits is a single point of technical, ethical, and legal failure.

ODDC requires ENVELO-compliant enforcement — an independent mechanism with three tiers of graduated response. Tier 1 monitors self-correction. Tier 2 forces controlled degradation to a Minimum Risk Condition. Tier 3 halts execution entirely. The enforcement layer does not evaluate intent or intelligence. It evaluates authority.

Q3: Won't the enforcement layer itself create new risks? Any safety mechanism introduces risk. The question is how it fails, not whether it can.

ENVELO-compliant enforcement uses a three-tier architecture designed to degrade gracefully. On ODD approach, the system self-corrects. On ODD breach, ENVELO forces a Minimum Risk Condition — the safest achievable state. Only when MRC fails does ENVELO halt execution entirely. If the enforcement mechanism itself cannot determine system state, it defaults to halt. Every tier transition is cryptographically logged.

Architectural isolation is a core requirement: the enforcement mechanism operates in a separate execution context with independent failure modes. A failure in the enforcement layer cannot propagate to the control model, and vice versa.

Technical Challenges Q4: Does the ENVELO Interlock add latency or degrade system performance?

The Interlock validates actions at execution time, which introduces a measurable but bounded overhead. The magnitude depends on the implementation architecture — hardware interlocks add nanoseconds, software interlocks in separate containers add microseconds to low milliseconds.

For context: financial trading systems already accept exchange-imposed latency from risk checks. Autonomous vehicles already run sensor fusion through redundant pipelines. The Interlock operates at a similar level — it is a validation gate, not a processing bottleneck.

ENVELO is a requirements specification, not software. Operators choose implementation approaches appropriate to their latency constraints. CAT-72 testing verifies that the chosen implementation meets enforcement requirements without compromising operational performance.

Q5: What if enforcement triggers incorrectly and causes harm — a false positive? A false positive — where the Interlock forces MRC or halt unnecessarily — is a real operational risk. The three-tier model is specifically designed to mitigate this.

Tier 2 (MRC) does not shut the system down. It transitions the system to the safest achievable state while preserving the ability to resume normal operation. A false Tier 2 event is an operational disruption, not a catastrophic failure. This is by design: the enforcement margin between the ODD boundary and the ENVELO wall provides a buffer zone for controlled response.

Tier 3 (hard halt) activates only when MRC is insufficient or the enforcement mechanism cannot determine system state. The threshold for Tier 3 is intentionally conservative — it defaults to safety. An unnecessary halt is preferable to an uncontrolled failure.

All enforcement events are cryptographically logged with full context. Post-incident analysis can distinguish false positives from legitimate enforcement, and operators can adjust their ODD boundaries or enforcement margins accordingly during re-certification.

Q6: What about AI systems managing other AI systems — recursive autonomy? Recursive autonomy — where one AI system orchestrates, monitors, or optimizes other AI systems — is explicitly within ODDC's scope. The Data Centers scenario in our industry applications document addresses this directly, including AI recursion depth as an ODD boundary parameter.

Each layer of the stack can have its own ODDC determination with its own ODD boundaries and enforcement tiers. The outer system's ODD includes constraints on the inner systems it manages — recursion depth limits, resource allocation ceilings, and cascade failure boundaries.

This is not a theoretical exercise. Hyperscale computing environments already run AI workload orchestrators managing AI inference clusters. ODDC provides the framework to enforce boundaries at each layer of that stack.

Q7: Is the adaptive learning period sufficient to discover all operational boundaries?

No learning period discovers all boundaries with certainty — the same way no test suite achieves complete coverage. The adaptive learning mode discovers boundaries from observed operational telemetry, which means it captures the system's actual behavioral envelope, not a theoretical specification.

The operator reviews and approves all auto-discovered boundaries before enforcement begins. This is a critical step: the operator brings domain expertise that telemetry alone cannot provide. They may tighten boundaries, add constraints the system hasn't encountered, or flag edge cases the learning period missed.

For high-risk domains or complex ODDs, Sentinel Authority may require extended test duration (up to 168 hours) or mandate the prescriptive path where the operator defines boundaries upfront based on engineering analysis, regulatory requirements, and manufacturer specifications. The two paths are complementary, not competing.

Gaming and Integrity Q8: What prevents operators from tuning behavior during the 72-hour test, then reverting afterward?

Three mechanisms address this directly. First, CAT-72 requires continuous, uninterrupted operation with cryptographic evidence generation. Any modification to system configuration, ODD parameters, or enforcement settings during the test period constitutes an interruption event requiring a full restart. The evidentiary chain is tamper-evident.

Second, ODDC is not a point-in-time certification. Phase 5 (Maintenance) requires ongoing ENVELO Interlock operation. The Interlock continues to monitor and enforce boundaries after certification. If the operator modifies the system in ways that violate the certified ODD, the Interlock detects the deviation.

Third, ODDC attests bounded operation at the time of determination. If the operator materially changes the system post-certification, the determination may no longer apply. This is analogous to a vehicle passing emissions testing and then having its catalytic converter removed — the test was legitimate, but the current state is non-conformant.

Ongoing Interlock monitoring catches this. Q9: What happens when the system is updated after certification?

Software updates are an expected part of autonomous system operation. ODDC addresses this through the maintenance framework.

Minor updates that do not alter the system's operational envelope or enforcement architecture are monitored through ongoing Interlock operation. If the updated system continues to operate within its certified ODD boundaries, the determination remains valid.

Material changes — new capabilities, expanded operational domains, modified decision logic that affects boundary behavior — require re-assessment. This may involve a new scope assessment and abbreviated or full CAT-72 re-testing depending on the nature of the change.

The operator is responsible for determining whether a change is material. However, the ongoing Interlock provides an objective signal: if the system begins approaching or breaching boundaries it previously stayed well within, that is evidence the operational envelope has changed.

Q10: Who certifies the certifier? What makes Sentinel Authority credible? This is the right question, and we take it seriously.

Sentinel Authority operates as an independent certification body — analogous to Underwriters Laboratories (UL) for electrical safety or Bureau Veritas for industrial inspection. We do not build, sell, or operate autonomous systems. We do not invest in companies that do. Our revenue comes from conformance determination services, which creates alignment with rigor, not leniency.

The ODDC framework itself is transparent: the requirements specifications, test procedures, and evidentiary standards are published documents. Any qualified party can evaluate whether a determination was conducted correctly by examining the cryptographic evidence chain.

We recognize that institutional credibility is earned, not declared. We actively engage with regulatory bodies, standards organizations, and the insurance industry to ensure ODDC reflects evolving best practices. We welcome external audit of our processes and determinations.

Regulatory and Legal Q11: How does ODDC interact with existing safety frameworks (NHTSA, FAA, FDA, IEC 61508)?

ODDC is complementary to existing safety frameworks, not competing with them. Safety certifications like IEC 61508, ISO 26262, DO-178C, and FDA's software validation guidance evaluate design processes, hazard analyses, and development lifecycle. They answer the question: "Was this system designed safely?"

ODDC evaluates a different question: "Does this system operate within defined boundaries with verifiable enforcement?" This is a runtime evidentiary question that existing frameworks do not address directly.

Regulators can use ODDC as a runtime enforcement layer that complements their existing oversight tools. An autonomous vehicle that passes NHTSA safety standards AND holds an ODDC determination provides both design-time and runtime assurance. Neither alone is sufficient; together they form a more complete picture.

Q12: Who is liable when enforcement activates? When it fails to? ODDC does not create, modify, or transfer liability. It creates evidence.

When enforcement activates correctly — the system approaches a boundary, self-corrects, degrades to MRC, or halts — the cryptographic audit trail provides contemporaneous evidence of what happened, when, and why. This evidence is useful to all parties: operators, regulators, insurers, and courts.

When enforcement fails — a boundary is breached without detection, an MRC is not achieved, or a halt does not execute — the evidentiary gap itself is informative. The absence of expected enforcement events in the audit trail is evidence of a conformance failure.

Liability allocation remains a matter of contract, regulation, and tort law. ODDC's contribution is making the factual record available. In a domain where incidents are often reconstructed from incomplete and disputed data, a tamper-evident enforcement log is a significant improvement.

Q13: Does ODDC work across jurisdictions? What about international operations? ODDC is jurisdiction-neutral by design. It attests operational facts — bounded operation with tiered enforcement — not regulatory compliance in any specific jurisdiction.

An ODDC determination is evidence that a system operates within defined boundaries with functioning enforcement. A regulator in the EU, the United States, Singapore, or Japan can evaluate that evidence against their own requirements and determine whether it satisfies their oversight needs.

This is deliberate. Regulatory frameworks vary significantly across jurisdictions, and they evolve at different rates.

ODDC provides a consistent evidentiary substrate that regulators can build upon, rather than attempting to harmonize regulatory requirements across borders — which is not our role.



Business and Adoption Q14: What happens when the AI needs to act outside its ODD in an emergency?

If the system acts outside its ODD, it is not ODDC-conformant for that action. This is the design, not a limitation.

ODDC attests bounded operation. Emergency behavior outside bounds is a different risk profile that requires different governance. The three-tier model provides graduated response precisely for these situations: Tier 2 (MRC) allows the system to transition to the safest achievable state rather than simply shutting down. The enforcement margin between the ODD boundary and the ENVELO wall is the controlled degradation window.

An autonomous vehicle's MRC is to pull to the shoulder. A surgical robot's MRC is to retract instruments and yield to the surgeon. These are not shutdowns — they are transitions to the safest achievable state given current context. The system remains operational; it is the autonomous authority that is constrained.

Q15: What's in it for operators? Why would anyone voluntarily adopt this? ODDC provides three forms of value that operators cannot create on their own.

First, verifiable evidence of operational discipline. Any operator can claim their system is safe. ODDC provides independent, third-party evidence — backed by a 72+ hour cryptographic audit trail — that the system actually operates within defined boundaries with functioning enforcement. This evidence is useful for regulatory engagement, insurance underwriting, and incident defense.

Second, standardization. Today, every operator invents their own way to describe operational boundaries and enforcement behavior. ODDC provides a common framework that regulators, insurers, and customers can evaluate consistently across vendors and industries.

Third, market differentiation. As regulatory scrutiny of autonomous systems increases, operators with ODDC attestation are positioned ahead of compliance requirements rather than scrambling to meet them. Early adopters set the standard that others will be measured against.

Q16: Do we have to share our proprietary algorithms or source code? No. Sentinel Authority certifies the process, not the parameters.

With the adaptive certification path (the default), the ENVELO Interlock observes your system in normal operation and auto-discovers operational boundaries from telemetry data. Your source code, model weights, decision logic, and proprietary algorithms are never accessed. The Interlock transmits only operational telemetry — what your system does, not how it decides.

You review and approve the auto-discovered boundaries before enforcement begins. Sentinel Authority verifies that enforcement works correctly; it never needs to know what's inside the box.

This is an intentional architectural decision, not a limitation. IP protection is a prerequisite for voluntary adoption. If conformance determination required source code access, the operators who most need independent oversight would be the last to seek it.

— End of Document —