



WHITE PAPER

The Accountability Chain

Why Autonomous System Liability Requires Independent Conformance Evidence

February 2026 — Public Document

Sentinel Authority

sentinelauthority.org

CONTENTS

1. The Multi-Party Problem

2. Case Law Analysis

- 2.1 Tesla Autopilot: Developer vs. Operator
- 2.2 Cruise Robotaxi: Developer vs. Deployer vs. Regulator
- 2.3 Tesla-Fanuc: Developer vs. Integrator vs. Deployer
- 2.4 Waymo: Developer vs. Software Version vs. Regulator
- 2.5 Healthcare AI: Developer vs. Deployer vs. Clinician

3. How ODDC Resolves Accountability

- 3.1 The Chain of Conformance
- 3.2 The Liability Shield
- 3.3 The Accountability Mirror

4. The Political Calculus

- 4.1 Why ODDC Cuts Both Ways
- 4.2 Historical Precedent

5. The Regulatory Opportunity

- 5.1 United States
- 5.2 European Union
- 5.3 The Global Convergence

6. Conclusion

Executive Summary

When an autonomous system causes harm, no one is accountable. Developers blame integrators. Integrators blame deployers. Deployers blame operators. Operators blame the software. This finger-pointing is not a failure of character—it is the structural consequence of deploying autonomous systems without independent conformance records that establish what each party was responsible for and whether they met that responsibility.

Autonomous systems involve multiple parties in a chain of development, integration, deployment, and operation. When these systems fail, determining accountability requires answering questions that current frameworks cannot resolve: Was the system operating within its designed parameters? Did the integrator configure it correctly? Did the deployer maintain it properly? Was the operator using it within approved conditions? Without independent records that answer these questions, every party in the chain has both the incentive and the plausible basis to deflect responsibility.

This white paper provides a comprehensive analysis of the accountability problem across autonomous vehicle, industrial robotics, and healthcare AI domains. It examines landmark cases and regulatory actions from 2023–2025 to demonstrate how the absence of independent conformance data produces predictable accountability failures. It then presents ODDC certification as the infrastructure that resolves accountability by establishing an independent, tamper-evident record of each party's conformance obligations and whether those obligations were met.

Critically, ODDC is politically bulletproof because it cuts both ways. For compliant actors, it provides concrete evidence that they met their obligations—compressing litigation, reducing insurance costs, and providing a liability shield. For non-compliant actors, it establishes precisely where the accountability chain broke. Neither industry nor regulators can object to a system that protects the responsible and exposes the irresponsible.

1. The Multi-Party Problem

Every autonomous system in production today involves at least four distinct parties with distinct responsibilities and distinct potential liability exposure:

The AI Developer builds the core autonomous algorithms, trains the models, and specifies the intended operating conditions. Liability theories against developers include design defect (the algorithm was inherently flawed), failure to warn (the developer knew of limitations and did not adequately communicate them), and negligent training (the model was trained on insufficient or biased data).

The System Integrator combines the AI software with hardware platforms, sensor suites, and mechanical systems. Integrators configure parameters, set safety thresholds, and validate that the complete system functions as intended. Liability theories include negligent integration, failure to validate system interactions, and inadequate safety analysis of the combined system.

The Deployer places the system into its operational environment. For autonomous vehicles, this is the fleet operator. For industrial robotics, this is the factory owner. For healthcare AI, this is the hospital or health system. Deployers are responsible for ensuring the system operates within its intended use conditions, maintaining the system, and training personnel. Liability theories include negligent deployment, failure to maintain, and operating outside approved conditions.

The Operator/User interacts with the system during operation. For partially autonomous vehicles, this is the driver. For industrial robots, this is the floor worker. For healthcare AI, this is the clinician. Liability exposure depends on the level of autonomy: as autonomy increases, operator liability typically decreases, but the transition points are legally undefined.

A RAND Corporation study of autonomous truck liability identified software companies, vehicle manufacturers, integrators, fleet operators, and even mapping data providers as potential liable parties in a single incident. A January 2026 academic study published in the World Electric Vehicle Journal found that four competing tort liability frameworks exist globally for autonomous vehicle accident governance, with no consensus on how to allocate responsibility across the supply chain.

In traditional product liability, the product either works as designed or it doesn't. With autonomous systems, the product may work exactly as designed by the developer, be correctly integrated by the integrator, be properly deployed by the deployer, and still cause harm because the operational conditions exceeded what any party anticipated. Without independent conformance records, determining which party—if any—failed their obligation is a multi-year, multi-million-dollar litigation exercise.

2. Case Law Analysis

Recent litigation provides concrete evidence of how the accountability gap produces inconsistent, expensive, and often unjust outcomes. The following cases illustrate the structural problem across domains.

2.1 Tesla Autopilot: Developer vs. Operator

The Florida jury's \$243 million verdict in the Tesla Autopilot fatality case allocated 33% liability to Tesla and 67% to the deceased driver. The jury found that Tesla had overstated the system's safety capabilities and had been unable to install adequate safety checks. This case establishes a critical precedent: the manufacturer can be held partially liable even when the human operator bears majority fault, if the manufacturer's representations about the system's capabilities were misleading.

But the case also illustrates the accountability problem. Tesla argued that the driver failed to pay attention despite explicit warnings. The plaintiff argued that Tesla's marketing materials and the system's name (Autopilot) created a false impression of capability. The jury split liability, but the allocation was based on contested expert testimony about what the system could and could not do—exactly the kind of evidence that an independent conformance record would have resolved

definitively. If Tesla's Autopilot had an independently verified ODD specification stating exactly what conditions it was designed to handle, and independently verified behavioral data showing whether it was operating within those conditions at the time of the crash, the liability allocation would have been a factual determination, not a battle of expert opinions.

2.2 Cruise Robotaxi: Developer vs. Deployer vs. Regulator

The Cruise robotaxi pedestrian incident in San Francisco produced a three-way accountability dispute. The pedestrian was initially struck by a human-driven vehicle and fell into the path of a Cruise robotaxi, which then dragged her approximately 20 feet before stopping. Cruise, as both developer and deployer, faced regulatory action from the California DMV and CPUC, and suspended all operations. Nearly 1,000 vehicles were recalled.

The accountability complexity extended beyond Cruise. The California DMV revoked Cruise's driverless testing permit after determining that the company had withheld information about the dragging portion of the incident. The CPUC suspended Cruise's commercial passenger service. General Motors, as Cruise's parent company, faced investor lawsuits alleging that safety concerns had been inadequately disclosed. The human driver of the vehicle that initially struck the pedestrian was also a party to the incident.

The case demonstrates how a single autonomous system incident can generate liability exposure across at least four parties—the developer/deployer (Cruise), the parent company (GM), the regulator (California DMV/CPUC for permitting operations), and the human driver. Without independent conformance records establishing what the robotaxi's ODD was, whether it was operating within that ODD, and whether the system's post-collision behavior was within specification, each party's liability required extensive adversarial discovery and expert analysis.

2.3 Tesla-Fanuc: Developer vs. Integrator vs. Deployer

The September 2025 lawsuit filed against both Tesla and Fanuc after a robotic arm struck a worker with approximately 8,000 pounds of force at Tesla's Giga Texas facility is the clearest illustration of the multi-party accountability problem. The injured worker's \$51 million suit names both the robot manufacturer (Fanuc) and the deployer/integrator (Tesla), alleging that the robot operated outside its programmed parameters.

The multi-defendant structure creates opposing defense theories. Fanuc's likely defense: the robot was correctly manufactured and the injury resulted from Tesla's integration and deployment decisions. Tesla's likely defense: the robot's behavior was a manufacturing or design defect that occurred during Fanuc's production process. Both parties have documentation supporting their positions. Neither party's documentation is independently verified. The litigation will take years and cost millions before a factual determination is reached about which party's responsibility was breached.

OSHA data adds statistical weight to this pattern: 77 robot-related workplace accidents between 2015 and 2022 produced 93 injuries including amputations, fractures, and crushing injuries. NIOSH documented 61 robot-related fatalities between 1992 and 2015. Korean occupational safety data shows 369 robot-related accidents in a single decade, averaging 27 to 49 incidents per year despite mandatory risk assessments. In each case, the accountability question—who was responsible for the

robot operating outside safe parameters?—depended on adversarial litigation rather than independent conformance evidence.

2.4 Waymo: Developer vs. Software Version vs. Regulator

Waymo's December 2025 recall of 3,067 robotaxis for school bus safety violations reveals a different accountability dimension: temporal accountability across software versions. The Austin Independent School District documented 20 incidents of Waymo vehicles illegally passing stopped school buses during the 2025–2026 school year. Five of these incidents occurred after Waymo had already deployed software updates that it believed would fix the problem.

The accountability question becomes: which software version caused each incident? If Waymo deployed a fix on November 17, 2025, and violations continued after that date, was the fix inadequate, or were those vehicles running older software? NHTSA's recall affected vehicles with production dates from August 20 to November 5, 2025. Without independent conformance records tracking which software version was running on which vehicle at which time, establishing temporal accountability requires forensic analysis of proprietary logs that may not have been designed for evidentiary use.

The Waymo case also raises the question of regulatory accountability. NHTSA's investigation was triggered by media reports and school district complaints, not by any automated detection system. If NHTSA had access to independent conformance monitoring data showing that Waymo vehicles were systematically failing to comply with traffic safety rules around school buses, the regulatory response could have been proactive rather than reactive. Children's safety should not depend on school bus camera footage going viral.

2.5 Healthcare AI: Developer vs. Deployer vs. Clinician

The 2024 analysis of 51 court cases involving software-related patient injuries spans three categories, each with distinct accountability patterns. Drug management systems that recommended dangerous dosages raised liability questions about the developer (who programmed the dosage logic), the deployer (the hospital that configured the system for its formulary), and the clinician (who followed the recommendation without independent verification). Clinical decision support tools that missed diagnoses raised questions about whether the developer's algorithm was at fault or whether the deployer failed to validate the tool against local patient populations. Surgical robotics injuries raised questions about whether the robot's behavior was a design defect, an integration failure, or operator error.

The American Medical Association has acknowledged that physician liability for AI-assisted decisions presents novel legal questions. A clinician who follows an AI recommendation that harms a patient can argue that the AI was supposed to be reliable. The developer can argue that the clinician was supposed to exercise independent clinical judgment. The hospital can argue that its governance policies (which may include ISO 42001 certification) addressed AI risk management. Each argument has merit. None has independent conformance evidence to resolve it.

3. How ODDC Resolves Accountability

3.1 The Chain of Conformance

ODDC resolves the accountability problem by establishing independent conformance obligations for each party in the autonomous system supply chain. Each party's certification addresses their specific responsibilities, creating a complete chain of accountability:

| Party | Without ODDC | With ODDC Certification |
|-------------------|---|--|
| AI Developer | Self-declares system capabilities in marketing materials and technical documentation. No independent verification of behavioral boundaries. Liability depends on adversarial expert analysis of proprietary code. | ODD formally specified in machine-readable format, independently verified. ENVELO Interlock enforcement requirements satisfied. CAT-72 behavioral testing confirms system operates within declared ODD. Conformance certificate provides courtroom-ready evidence of what the developer committed to and whether the system meets that commitment. |
| System Integrator | Integration documentation exists but is not independently verified against developer's specifications. No standardized validation that integration preserved safety boundaries. Liability depends on comparing proprietary documentation from two parties. | Integration conformance verified against developer's certified ODD. ENVELO Interlock integrity confirmed post-integration. CAT-72 testing repeated on integrated system. Conformance certificate establishes that the integrator preserved the developer's certified behavioral boundaries. |
| Deployer | Deployment conditions self-assessed against manufacturer guidance. No independent verification that operating environment matches ODD. Maintenance records may not include software version tracking. Liability depends on reconstructing operational conditions from incomplete records. | Deployment conformance verified against certified ODD. Operating environment validated against formal boundary specifications. Continuous conformance monitoring with tamper-evident records. Conformance certificate establishes that the deployer maintains the system within its certified operating conditions. |
| Insurer | Underwrites based on self-reported data with no independent verification. Cannot assess correlated risk across fleet. Cannot differentiate between compliant and non-compliant operators. Pricing based on industry averages, not individual conformance. | Underwrites based on independently verified conformance status. Can assess correlated risk through ODD + Interlock version tracking. Can offer risk-differentiated pricing based on conformance. Conformance data provides actuarial-grade evidence for loss modeling. |

3.2 The Liability Shield

For compliant actors, ODDC certification provides a concrete, independently verified liability defense. When a certified system is involved in an incident:

For the developer: The ODDC conformance certificate establishes exactly what behavioral boundaries the system was designed to operate within and provides independently verified evidence that the system met those boundaries in testing. If the incident occurred within the certified ODD and the system's behavior was consistent with its certified specifications, the developer has concrete evidence that the product was not defective. If the incident occurred outside the ODD, the conformance record establishes that the system was being operated beyond its certified conditions—shifting accountability to the deployer or operator.

For the integrator: The conformance certificate for the integrated system establishes that the integration preserved the developer's certified behavioral boundaries. If post-integration testing confirmed behavioral conformance and the ENVELO Interlock integrity was verified, the integrator has evidence that their work did not introduce the failure.

For the deployer: The continuous conformance monitoring record establishes that the system was being operated within its certified ODD at the time of the incident. If conformance was maintained, the deployer has evidence of responsible operation. If conformance had lapsed, the deployer's exposure is defined, but it is bounded and documented rather than open-ended.

3.3 The Accountability Mirror

For non-compliant actors, ODDC creates the inverse effect: the absence of independent conformance evidence becomes itself an accountability factor. When an uncertified system is involved in an incident:

The developer cannot demonstrate that the system was designed to operate within specific behavioral boundaries, because no independently verified boundaries exist. The integrator cannot demonstrate that the integration preserved safety properties, because no baseline was independently established. The deployer cannot demonstrate that the system was operated within its design parameters, because no independent monitoring data exists. Each party's defense depends on their own proprietary records, which the adversary will challenge as self-serving.

This asymmetry—the advantage that certification provides to compliant actors and the disadvantage that its absence creates for non-compliant actors—produces a natural market incentive for certification. Over time, the standard of care for deploying autonomous systems in safety-critical applications will include independent conformance verification, just as the standard of care for selling electrical products includes UL certification and the standard of care for deploying aircraft includes FAA airworthiness certification.

4. The Political Calculus

4.1 Why ODDC Cuts Both Ways

ODDC is designed to be politically bulletproof because it serves both sides of every policy debate in autonomous systems governance:

For industry advocates who argue that regulation stifles innovation: ODDC is not a regulatory mandate. It is an independent certification that provides competitive advantage to companies that invest in safety. Certified companies get lower insurance costs, faster regulatory approvals, and concrete liability defenses. ODDC rewards innovation in safety rather than penalizing innovation in capability.

For safety advocates who argue that industry self-regulation is insufficient: ODDC provides independent, third-party verification that cannot be gamed by the entity being verified. The ENVELO Interlock is non-bypassable by design. CAT-72 testing is conducted under observed conditions with tamper-evident results. ODDC addresses the specific concern that self-reported safety data is unreliable.

For Republican legislators focused on reducing regulatory burden: ODDC is a market-based mechanism that achieves safety outcomes without new regulation. It provides the evidentiary infrastructure that allows existing liability law to function effectively for autonomous systems, reducing the need for prescriptive rules.

For Democratic legislators focused on consumer protection: ODDC ensures that autonomous system manufacturers cannot externalize safety costs by deploying uncertified systems. It provides consumers and communities with independently verified evidence of system safety, rather than relying on corporate self-attestation.

4.2 Historical Precedent

The politically bulletproof nature of independent certification is not theoretical—it has been demonstrated repeatedly in the history of safety certification:

Underwriters Laboratories (UL): Founded in 1894 by insurers who recognized that self-reported electrical safety claims were unreliable, UL became the de facto standard for consumer product safety without a government mandate. Manufacturers adopted UL certification because retailers required it, insurers priced it in, and consumers trusted it. No legislation mandates UL certification—but try selling an electrical product without it.

Lloyd's Register: Marine insurance created independent vessel classification because underwriters could not price hull risk without independent structural verification. Today, classification society certification is required by international maritime law, but it began as a market-driven mechanism.

TÜV (Technischer Überwachungsverein): German technical inspection associations emerged in the 19th century to inspect steam boilers—a technology whose failure could be catastrophic. TÜV certification became the standard because operators who invested in safety deserved to be distinguished from those who did not.

In each case, the certification body succeeded because its existence served everyone's interests: manufacturers who invested in safety gained competitive advantage; insurers gained pricing data; regulators gained enforcement evidence; and the public gained protection. ODDC follows this same model for autonomous systems.

5. The Regulatory Opportunity

5.1 United States

The current U.S. regulatory environment is fragmented and reactive. NHTSA's enforcement authority is limited to post-hoc recall after failures have already occurred. The bipartisan Autonomous Vehicle Acceleration Act (S. 1798) addresses FMVSS modernization and state preemption but does not establish independent conformance requirements. State-level regulations vary dramatically, with Kentucky raising AV truck insurance minimums to \$5 million while other states impose no specific AV insurance requirements.

ODDC provides a mechanism that works within this fragmented environment. Federal regulators can reference ODDC conformance in guidance documents without new rulemaking. State regulators can incorporate ODDC requirements into insurance mandates. Congressional committees can cite ODDC as evidence that the private sector is developing solutions that reduce the need for prescriptive regulation. The fact that ODDC already provides an enforcement infrastructure that regulators can leverage—rather than one they must build—makes it attractive to agencies facing budget constraints and political opposition to new mandates.

5.2 European Union

The EU AI Act's conformity assessment requirements create a direct market for independent behavioral verification. With harmonized standards delayed and the first standard (prEN 18286) still in public enquiry, providers of high-risk AI systems face an August 2026 compliance deadline without clear guidance on how to generate the required conformity evidence. ODDC provides a deployable conformance methodology that can serve as a bridge while harmonized standards mature.

For EU policymakers, ODDC addresses a specific concern: the risk that conformity assessment requirements become a paper exercise. If conformity assessment is reduced to management system certification (ISO 42001) without behavioral verification, the EU AI Act's safety objectives will not be met. ODDC provides the behavioral layer that gives conformity assessment teeth.

5.3 The Global Convergence

The January 2026 academic study comparing tort liability frameworks across jurisdictions found four competing models globally. ODDC provides a standardized evidentiary framework that works across all four models. Regardless of whether a jurisdiction applies strict liability, negligence, no-fault, or comparative fault principles, independent conformance evidence establishes the same factual baseline:

what the system was designed to do, whether enforcement mechanisms were in place, and whether the system was operating within its declared boundaries at the time of an incident.

6. Conclusion

The accountability chain in autonomous systems is broken because no independent mechanism exists to establish what each party was responsible for and whether they met that responsibility. This produces predictable consequences: prolonged litigation, inconsistent outcomes, misallocated liability, and perverse incentives that reward opacity over transparency.

ODDC repairs the chain by establishing independent conformance obligations for each party—developer, integrator, deployer—with tamper-evident records that provide courtroom-ready evidence of compliance or non-compliance. For compliant actors, this evidence is a shield. For non-compliant actors, its absence is an exposure. The system rewards investment in safety and penalizes neglect of it.

The accountability chain does not need new liability law. It needs independent evidence. ODDC provides that evidence. The result is a system where doing the right thing is rewarded, doing the wrong thing is exposed, and neither side of the political spectrum can object—because it protects the responsible and holds the irresponsible to account.

REFERENCES

- [1] Florida Jury Verdict, Tesla Autopilot Fatal Accident, Key Largo. \$243 million award, 33% manufacturer liability, 67% operator. 2025.
- [2] Cruise Robotaxi Pedestrian Incident, San Francisco. California DMV permit revocation, CPUC service suspension, ~1,000 vehicle recall. 2023.
- [3] Tesla-Fanuc \$51M Lawsuit. Robotic arm struck worker with ~8,000 lbs force at Giga Texas, Waller County, TX. September 2025.
- [4] Waymo Recall of 3,067 Robotaxis. NHTSA investigation into school bus safety violations. December 2025. 20 incidents documented by Austin ISD.
- [5] RAND Corporation. "Liability for Autonomous Truck Accidents." Identifies software companies, vehicle manufacturers, integrators, fleet operators, and mapping providers as potential liable parties.
- [6] Long, B., Zhao, Z., and Cai, Q. "Comparing Tort Liability Frameworks in Autonomous Vehicle Accident Governance." World Electric Vehicle Journal, 17(1):32, January 2026. Four competing global frameworks.
- [7] OSHA. Robot-Related Workplace Accidents: 77 incidents resulting in 93 injuries including amputations, fractures, crushing (2015–2022).
- [8] NIOSH. Robot-Related Fatalities in the United States: 61 deaths (1992–2015).
- [9] Korean Ministry of Employment and Labor. 369 robot-related accidents in a single decade (27–49/year).
- [10] Bates, D. et al. "Software-Related Patient Injuries: Analysis of 51 Court Cases." Drug management, clinical decision support, surgical robotics. 2024.
- [11] American Medical Association. Statement on physician liability for AI-assisted clinical decisions. Novel legal questions acknowledged.
- [12] NHTSA. Automated Vehicle Framework, updated April 2025. Standing General Order on ADS incident reporting.

- [13] Autonomous Vehicle Acceleration Act of 2025 (S. 1798). FMVSS modernization and state preemption.
- [14] Kentucky SB 241 (2025). AV truck insurance minimum raised to \$5M; human-on-board extended to 2031.
- [15] European Union. Regulation (EU) 2024/1689 (AI Act). High-risk conformity assessment requirements.
- [16] ECRI. "Top 10 Health Technology Hazards for 2025." AI without proper oversight ranked #1.
- [17] Stanford HAI. "Assessment of Trustworthy AI in the Context of Healthcare." Tools tested only by developers.
- [18] Austin Independent School District. Letter to Waymo: 20 illegal school bus passes, including post-fix incidents. 2025–2026.
- [19] Swiss Re and Waymo. "Comparative Safety Performance of Autonomous- and Human Drivers." December 2024.
- [20] Sentinel Authority. ODDC Overview v3.0, ENVELO Requirements v3.0, CAT-72 Procedure v3.0. Published at sentinelauthority.org.

This document references publicly available government reports, court filings, regulatory publications, and peer-reviewed research. This document does not constitute legal advice.